

CHROMATIN ARCHITECTURE ABERRATIONS IN PROSTATE CANCER AND LEUKEMIA

by

James Hawley

A thesis submitted in conformity with the requirements
for the degree of Doctor of Philosophy

Graduate Department of Medical Biophysics
University of Toronto

© Copyright 2021 by James Hawley

Contents

1 Hedging uncertainty in differential gene expression analyses with James-Stein estimators	1
1.1 Abstract	1
1.2 Introduction	1
1.3 Results	2
1.3.1 Derivation of the James-Stein fold change estimator	2
1.3.2 Comparison between the OLS and James-Stein estimators	4
1.3.3 Empirical analysis of the James-Stein estimator	6
1.4 Discussion	9
1.5 Methods	9
1.5.1 RNA sequencing data collection and pre-processing	9
1.5.2 Differential expression analysis	9
1.5.3 Random sampling and simulations	9
1.5.4 Random sampling of smaller numbers of transcripts	10
A Supplementary Material for Chapter 4	11
A.1 Differential expression analysis with Sleuth	11
A.2 Statistical moments of the ordinary least squares estimator	13
A.3 Statistical moments of the James-Stein estimator	14
A.3.1 Expected value of the James-Stein estimator	14
A.3.2 Variance of the James-Stein estimator	15
Glossary	17
References	19

Chapter 1

Hedging uncertainty in differential gene expression analyses with James-Stein estimators

J.R.H., and M.L. conceptualized the study. J.R.H. derived the statistical estimates and designed and conducted all the experiments. Figures were designed by J.R.H. The manuscript was written by J.H., and M.L. M.L. oversaw the study.

1.1 Abstract

1.2 Introduction

- the two main approaches for reducing error in a model are to reduce the model variance or model bias Figure 1.1
- here we attempt to decrease mean square error (MSE) by simultaneously increasing the bias and decreasing the variance in fold change coefficient estimators
- derivation for the James-Stein (JS) estimator can be found in Section 1.3.1
- Equation for the JS estimator can be seen in Figure 1.1b.
- in theory, this may increase the error of some transcripts, but will decrease MSE for a set of transcripts in aggregate Appendix A.3

First, we derive the JS estimator fold gene expression fold change and relate it to the ordinary least squares (OLS) estimator. Then, using simulations from a highly replicated RNA sequencing (RNA-seq) experiment [1], we compare the differences in statistical inferences between the JS and OLS estimators. Finally, we investigate how the number of transcripts under consideration affects the reduction in MSE, suggesting how this method can be best used in practice.

1.3 Results

1.3.1 Derivation of the James-Stein fold change estimator

For a p -variate normal distribution $Z \sim \mathcal{N}_p(\mu, \Sigma)$ where μ is unknown and Σ is known, the following theorem holds [2]:

Theorem 1 *The estimator $\hat{\mu}^{(0)} = Z$, for any mean μ , does not minimize the MSE $\mathbb{E}[(\mu - \hat{\mu})^2]$ for a single observation of Z when $p \geq 3$ and $\Sigma = I_p$. Namely, the estimator $\hat{\mu}^{(JS)} = \left(1 - \frac{b}{a + \|Z\|^2}\right) Z$ has a smaller MSE than $\hat{\mu}^{(0)}$ for a single observation for sufficiently small b and large a .*

This result was generalized to non-singular covariance matrices that were not necessarily the identity matrix (Theorem 2 of [REF 3]):

Theorem 2 *Let $Z \sim \mathcal{N}_p(\mu, \Sigma)$. Let $\hat{\mu}^{(JS)} = \left(1 - \frac{c}{Z^T \Sigma^{-1} Z}\right) Z$. If $p \geq 3$, $\text{Tr}(\Sigma) \geq 2\lambda_L$ (where λ_L is the largest eigenvalue of the covariance matrix, Σ), and $0 \leq c \leq 2 \left(\frac{\text{Tr}(\Sigma)}{\lambda_L} - 2\right)$, then $\hat{\mu}^{(JS)}$ is the estimator for the mean, μ , that minimizes the MSE for a single observation of Z .*

Consider the differential analysis model used in the Sleuth R package [4, 5] for a single transcript, s , with the simple experimental design of n_{wt} wild-type (WT) samples and 1 mutant sample:

$$D_s | Y_s \sim \mathcal{N}_N(\beta_{s,0} + \mathbb{I}_{mut} \beta_{s,1}, (\sigma_s^2 + \tau_s^2) I_N)$$

For the n_{wt} WT samples, this is equivalent to

$$D_s | Y_s \sim \mathcal{N}_{n_{wt}}(\beta_{s,0}, (\sigma_s^2 + \tau_s^2) I_{n_{wt}})$$

which can be fit with the same model process that Sleuth employs. For the single mutated sample, this model is

$$D_s | Y_s \sim \mathcal{N}(\beta_{s,0} + \beta_{s,1}, \max\{\hat{\sigma}_s^2, \tilde{\sigma}_s^2\} + \hat{\tau}_s^2) \quad (1.1)$$

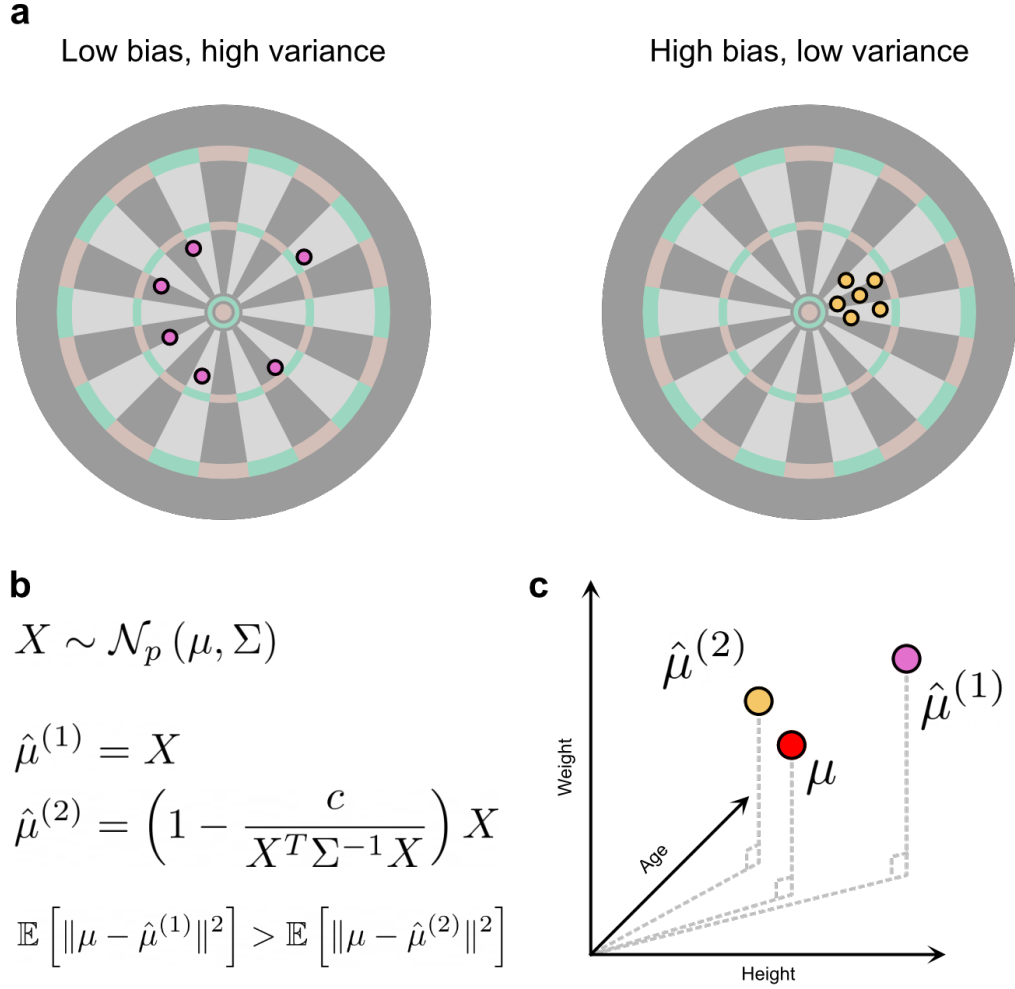


Figure 1.1: **Reducing the bias-variance tradeoff by combining information across multiple features.** **a.** Schematic of the bias-variance tradeoff for assessing model performance. Dartboard on the left shows low bias of the darts (mean is close to the bullseye) but a large variance. Dartboard on the right shows a high bias of the darts (mean is off-centre), but a small variance. **b.** For a p -variate normal distribution from which a single observation is made, the naive estimator has a higher MSE than the JS estimator, defined as $\hat{\mu}^{(2)}$. **c.** An analogy showing how the JS estimators work in theory. Trying to estimate the mean height, weight, and age for the entire population (μ) from a single person will give an estimate that is likely far from the truth ($\hat{\mu}^{(1)}$). Combining information from the three variables together can produce an estimate that is closer to the truth ($\hat{\mu}^{(2)}$).

The covariance matrix is the same as the mutated samples, but the mean $\beta_{s,0} + \beta_{s,1}$ is unknown. By reparameterizing this model to consider every transcript in the single mutated sample, Equation (1.1) can be re-written as follows:

$$\Delta \sim \mathcal{N}_{|S|}(\mathbf{B}_0 + \mathbf{B}_1, \Sigma) \quad (1.2)$$

where

$$\begin{aligned} \mathbf{B}_{i,s} &= \beta_{s,i} \forall s \in S \\ \Sigma &= \begin{bmatrix} \max\{\hat{\sigma}_1^2, \tilde{\sigma}_1^2\} + \hat{\tau}_1^2 & & 0 \\ & \ddots & \\ 0 & & \max\{\hat{\sigma}_{|S|}^2, \tilde{\sigma}_{|S|}^2\} + \hat{\tau}_{|S|}^2 \end{bmatrix} \end{aligned}$$

We switch from using coefficients $\beta_{t,i}$ to $\mathbf{B}_{i,s}$ to avoid confusion, since $\beta_{t,i} \in \mathbb{R}^p$ (a p -dimensional vector for each covariate in the design) whereas $\mathbf{B}_{i,s} \in \mathbb{R}^{|S|}$ (an $|S|$ -dimensional vector for only a single coefficient over all transcripts in S).

Observations of a single mutated sample from this model meet the criteria for the JS estimators. A JS estimator for the unknown fold change, \mathbf{B}_1 , can be constructed.

$$\hat{\mathbf{B}}_1^{(JS)} = \left(1 - \frac{c}{(\Delta - \hat{\mathbf{B}}_0)^T \Sigma^{-1} (\Delta - \hat{\mathbf{B}}_0)} \right) (\Delta - \hat{\mathbf{B}}_0) \quad (1.3)$$

where $\hat{\mathbf{B}}_0$ is the estimate obtained from the non-mutated samples for all transcripts $s \in S$.

It is straightforward to see that $\text{Tr}(\Sigma) = \sum_{s \in S} \max\{\hat{\sigma}_s^2, \tilde{\sigma}_s^2\} + \hat{\tau}_s^2$ and that $\lambda_L = \max_{s \in S} \{\max\{\hat{\sigma}_s^2, \tilde{\sigma}_s^2\} + \hat{\tau}_s^2\}$.

1.3.2 Comparison between the OLS and James-Stein estimators

For a simple experimental design where the mutation status is the only coefficient the OLS estimator is given by:

$$\begin{bmatrix} \hat{\beta}_{s,0}^{(OLS)} \\ \hat{\beta}_{s,1}^{(OLS)} \end{bmatrix} = \hat{\beta}_s^{(OLS)} = (X^T X)^{-1} X^T d_s = \begin{bmatrix} \bar{d}_s^{(wt)} \\ d_s^{(mut)} - \bar{d}_s^{(wt)} \end{bmatrix}$$

where

$$X = \begin{bmatrix} 1 & 1 \\ 1 & 0 \\ \vdots & \vdots \\ 1 & 0 \end{bmatrix} \in \mathbb{R}^{(N+1) \times 2}$$

is the design matrix. Looking closely at the OLS estimator for the mutation coefficient, $\beta_{s,1}$, it is clear that it is given by:

$$\hat{\beta}_{s,1}^{(OLS)} = d_s^{(mut)} - \hat{\beta}_{s,0}^{(OLS)} = \delta_s - \hat{\beta}_{0,s} \quad (1.4)$$

which is used directly in the definition of the JS estimator in Equation (1.3). The JS estimator for B_1 can then be expressed simply as:

$$\hat{B}_1^{(JS)} = \left(1 - \frac{c}{\left(\hat{B}_1^{(OLS)} \right)^T \Sigma^{-1} \hat{B}_1^{(OLS)}} \right) \hat{B}_1^{(OLS)} \quad (1.5)$$

From this definition, it is easy to see that the JS estimate is colinear with the OLS estimate but uniformly shrunk towards 0. For a more general experimental design, the above can be extended.

Theorem 3 *Given an experimental design matrix $X \in \mathbb{R}^{n \times p}$ where $n > p$, $\text{rank}(X) = p$ and $\text{rank}(X^*) = p - 1$ where $X^* \in \mathbb{R}^{(n-1) \times p}$ is the same design matrix but with one sample removed, a JS estimator for the linear coefficient uniquely specified by the one sample is given by*

$$\hat{B}_i^{(JS)} = \left(1 - \frac{c}{\left(\hat{B}_i^{(OLS)} \right)^T \Sigma^{-1} \hat{B}_i^{(OLS)}} \right) \hat{B}_i^{(OLS)}$$

It can be shown that the JS estimator is biased towards 0 with a smaller variance than the OLS estimator (see Appendix A.3).

From this definition, we can see two parameters of this model that will affect the amount of biasing: the scaling coefficient, c , and the total number of transcripts being considered, $|S|$. Firstly, the scaling coefficient can be manually specified, and the largest biasing occurs when c is its maximum value, $2 \left(\frac{\text{Tr}(\Sigma)}{\lambda_L} - 2 \right)$. Secondly, the transcripts under consideration can also be manually specified, which will affect the value of the denominator $\left(\hat{B}_1^{(OLS)} \right)^T \Sigma^{-1} \hat{B}_1^{(OLS)}$, and thus the amount of biasing. The more transcripts under consideration, the larger the denominator, and thus the smaller the effect, compared to the OLS method. Thus, we have produced a high-bias, low-variance fold change estimator that has a lower MSE than the OLS estimator and two tunable parameters.

1.3.3 Empirical analysis of the James-Stein estimator

Now that the statistical properties of the JS estimator have been shown, we use empirical RNA-seq data to measure the performance of this method in practice. To demonstrate this, we make use of a highly replicated RNA-seq experiment involving Δ Snf2 knockout (KO) and WT yeast cells [1]. This dataset contains 48 biological replicates of each condition, an infeasible sample size for most RNA-seq experiments. Experiments with small sample sizes can be compared to the full dataset to estimate the number of true and false detections for a given experimental design and a given method. We randomly select N total samples from the full dataset with an optimal even split between the two groups (i.e. $N/2$ Δ Snf2 and $N/2$ WT) or a suboptimal $N - 1$ -vs-1 split (i.e. $N - 1$ Δ Snf2 and 1 WT or 1 Δ Snf2 and $N - 1$ WT). To measure the effect of total sample size, these simulations are repeated for multiple values of N .

We find that for all values of N , the JS method produces more true positive (TP) and false positive (FP) calls, as well as fewer true negative (TN) and false negative (FN) calls, than the OLS method with suboptimal designs, on average (two-way $??$, $p < 2.2 \times 10^{-16}$; Figure 1.2). For example, for $N = 6$, the JS method identified 642.87 TP, 58 FP, 2098.3 TN, and 2995.7 FN on average, compared the to 520.7 TP, 41.37 FP, 2114.97 TN, and 3117.93 FN for the OLS method (Figure 1.2). The strength of this effect appears to decrease as the total number of samples increases. Notably, for the $N = 4$ case where a suboptimal design would be most common in practice, the JS method had more TP and fewer FN than the optimal experimental design. In all other cases, however, the optimal even split between Δ Snf2 and WT groups results in the highest TP and fewest FN calls, as expected. Thus, for differential expression hypothesis testing, the OLS method can identify more TP and fewer FN calls than the OLS method when dealing with suboptimal experimental designs.

From the derivation of the JS estimator, it

- to investigate where these changes in fold change estimates and p -values originates from, we investigate representative simulations
- most genes in the Δ Snf2 model are under-expressed compared to the WT model Figure 1.3

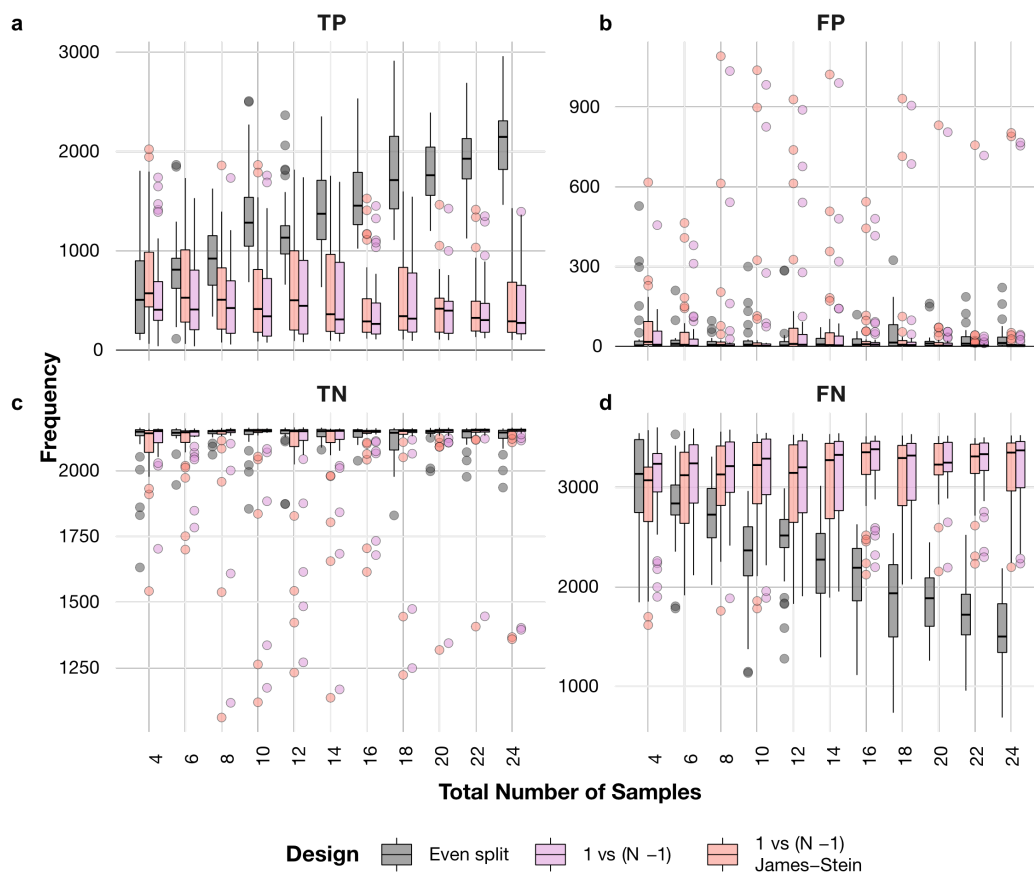


Figure 1.2: **Differential gene expression analysis of the entire yeast transcriptome with differently sized experimental designs.** Simulations ($n = 30$) using randomly selected samples which were then compared to the full dataset of 48 Δ Snf2 vs 48 WT to calculate TP (a), FP (b), TN (c), and FN (d).




Figure 1.3: **Differential gene expression analysis of Δ Snf2 vs WT yeast cells using different sample sizes and experimental designs.** **a.** Volcano plot of differential expression results with OLS estimates in a highly replicated experiment consisting of 48 biological replicates of each condition. **b.** The same analysis as (a) using 4 samples in total, 2 Δ Snf2 and 2 WT samples. **c.** The same analysis as (c) using 1 Δ Snf2 and 3 WT. **d.** The same analysis as (c) using the JS method instead of OLS.




Figure 1.4: **Differential gene expression analysis focusing on a subset of transcripts, not the entire transcriptome.** All experiments use 1 Δ Snf2 vs 5 WT samples (or vice versa). **a.** Comparison of the MSE of the JS estimates (y -axis) against the OLS estimates (x -axis). The total number of transcripts in each comparison is specified above each facet. **b.** Percent different in MSE between the JS and OLS estimates. One-sided, two-sample Student's t -test, $n = 30$, FDR multiple test corrections.

1.4 Discussion

1.5 Methods

1.5.1 RNA sequencing data collection and pre-processing

RNA-seq reads were obtained from [REF 1] from the Sequence Read Archive (SRA) (Project Accession PRJEB5348). Briefly, the Δ Snf2 KO and WT cells were both sequenced as single-end, 50 bp reads, split across seven lanes of a single Illumina HiSeq 2000 flow cell. The quality of the sequencing reads was assessed with FastQC [6]. The *Saccharomyces cerevisiae* R64-1-1 reference transcriptome index from Ensembl (v96) was downloaded from <https://github.com/pachterlab/kallisto-transcriptome-indices/>. RNA-seq reads from all technical replicates of a single biological replicate were quantified against the reference transcriptome with Kallisto (v0.46.2) [7] with the following command:

```
kallisto quant --bootstrap-samples 100 --pseudobam --threads 8 --
    index /path/to/R64-1-1.idx --output-dir {output_dir} {input.fastq
    .gz} > {output_report}
```

1.5.2 Differential expression analysis

Differential expression analysis was conducted in R (v4.0.2) [8] with the Sleuth package (v0.30.0) [4, 5]. Statistical significance for the fold change of transcripts from the mutation status was determined with a two-sided Wald test. Multiple testing correction was performed using the Benjamini-Hochberg FDR method [9]. Transcripts were determined to be significantly differentially expressed if $FDR < 0.01$.

This method was used for differential expression analysis for the full experimental design (48 Δ Snf2 replicates and 48 WT replicates), as well as all simulations with smaller sample sizes.

1.5.3 Random sampling and simulations

From the 48 biological replicates of the Δ Snf2 and WT samples, N total samples were randomly selected, where $N \in \{4, 6, 8, \dots, 22, 24\}$. The sampling was repeated 30 times, independently. The samples were chosen according to of the following experimental designs:

1. 30 iterations of $N/2$ Δ Snf2 vs $N/2$ WT, compared using the OLS estimator

2. 15 iterations of 1 Δ Snf2 vs $N - 1$ WT + 15 iterations of $N - 1$ Δ Snf2 vs 1 WT, compared using the OLS estimator
3. 15 iterations of 1 Δ Snf2 vs $N - 1$ WT + 15 iterations of $N - 1$ Δ Snf2 vs 1 WT, compared using the JS estimator

For each simulation, the number of TP, FP, TN, and FN calls was calculated by comparing the classification of a given transcript as significantly differentially expressed or not (i.e. if $FDR < 0.01$ for a given transcript) in the simulation to the full experiment of 48 Δ Snf2 vs 48 WT. With this confusion matrix, derived rates, such as the true positive rate, were then calculated.

1.5.4 Random sampling of smaller numbers of transcripts

In a similar fashion to above, 30 iterations of randomly selected samples were chosen to match the same three scenarios as above, with a fixed $N = 6$ (15 iterations with 1 Δ Snf2 vs 5 WT and 15 iterations with 5 Δ Snf2 vs 1 WT). A subset of transcripts, S , were randomly selected from those transcripts with ≥ 10 estimated reads in all 6 of the randomly selected samples. Fold change estimates for the three scenarios were calculated as above. The number of transcripts selected was chosen from $|S| \in \{3, 10, 25, 50, 100, 250\}$.

MSE was calculated by comparing the fold change estimate from each iteration to the fold change estimate from the full experiment with 48 Δ Snf2 vs 48 WT. Percent differences in MSE were calculated as:

$$\text{Percent difference} = \frac{MSE_{JS} - MSE_{OLS}}{MSE_{OLS}}$$

Appendix A

Supplementary Material for Chapter 4

A.1 Differential expression analysis with Sleuth

The differential expression model employed in the Sleuth (v0.30.0) [4, 5] can be described as follows. Consider a set of transcripts, S , measured in N samples with an experimental design matrix, $X \in \mathbb{R}^{N \times p}$, where p is the number of covariates considered. Let Y_{si} be the natural log of the abundance of transcript s in sample i . Given the design matrix

$$X = [x_1^T; x_2^T; \dots x_n^T], x_i \in \mathbb{R}^p$$

the abundance of transcripts can be modelled as a generalized linear model (GLM)

$$Y_{si} = x_i^T \beta_s + \epsilon_{si} \tag{A.1}$$

where $\epsilon_{si} \sim \mathcal{N}(0, \sigma_s^2)$ is the biological noise of transcript s in sample i and $\beta_s \in \mathbb{R}^p$ is the fixed effect of the covariates on the expression of transcript s .

Due to inferential noise from sequencing, each Y_{si} are not observed directly, but indirectly through the observed perturbations, D_{si} . This can be modelled as

$$D_{si}|Y_{si} = Y_{si} + \zeta_{si} \tag{A.2}$$

where $\zeta_{si} \sim \mathcal{N}(0, \tau_s^2)$ is the inferential noise of transcript s in sample i . Both biological and inferential noise for each transcript are independent and identically distributed (IID) and independent of each other. Namely:

$$\mathbb{C}ov[\epsilon_{si}, \epsilon_{rj}] = \sigma_s^2 \delta_{i,j} \delta_{s,r}$$

$$\mathbb{C}ov[\zeta_{si}, \zeta_{rj}] = \tau_s^2 \delta_{i,j} \delta_{s,r}$$

$$\mathbb{C}ov[\epsilon_{si}, \zeta_{rj}] = 0$$

$$\forall s, r \forall i, j$$

The abundances for transcript s in all N samples can then modelled as a multivariate normal distribution

$$D_s | Y_s \sim \mathcal{N}(X\beta_s, (\sigma_s^2 + \tau_s^2)I_N) \quad (\text{A.3})$$

where $I_N \in \mathbb{R}^{N \times N}$ is the identity matrix.

The goal of the differential analysis is to estimate the $|S| \times p$ coefficients in $B_s \forall s \in S$, and to determine which coefficients differ significantly from 0. This is achieved through a Wald test or likelihood ratio test after estimating the inferential variance, τ_s^2 , through bootstrapping and the biological variance, σ_s^2 , through dispersion estimation and shrinkage.

The estimator for the differential effect is the OLS estimate:

$$\hat{\beta}_s = (X^T X)^{-1} X^T d_s$$

where d_s is the observed abundances given by

$$d_{si} = \ln \left(\frac{k_{si}}{\hat{f}_i} + 0.5 \right)$$

$$\hat{f}_i = \text{median}_{s \in S^*} \frac{k_{si}}{\sqrt[N]{\prod_{j=1}^N k_{sj}}}$$

where k_{si} is the estimated read count from the Kallisto package (v0.46.1) [7] for transcript s in

sample i and \hat{f}_i is the scaling factor for sample i , calculated from the set of all transcripts that pass initial filtering, S^* .

A.2 Statistical moments of the ordinary least squares estimator

As shown in Supplementary Note 2 of [REF 4], the estimator is unbiased, Namely

$$\mathbb{E} \left[\hat{\beta}_s^{(OLS)} \right] = B_s \quad (\text{A.4})$$

It can also be shown that, for a covariance matrix Σ ,

$$\mathbb{V} \left[\hat{\beta}_s^{(OLS)} \right] = (X^T X)^{-1} X^T \Sigma X (X^T X)^{-1}$$

In the case where $\Sigma = (\sigma_s^2 + \tau_s^2) I_N$, this reduces to

$$\mathbb{V} \left[\hat{\beta}_s^{(OLS)} \right] = (\sigma_s^2 + \tau_s^2) (X^T X)^{-1}$$

Consider a simple experimental design where the only covariate of interest is the presence of a mutation. Then the design matrix, with the first column being the intercept and the second being the mutation status, looks like so:

$$X = \begin{bmatrix} 1 & 1 \\ 1 & 0 \\ \vdots & \vdots \\ 1 & 0 \end{bmatrix} \in \mathbb{R}^{(N+1) \times 2}$$

The variance of the OLS estimator is then

$$\mathbb{V} \left[\hat{\beta}_s^{(OLS)} \right] = \frac{(\sigma_s^2 + \tau_s^2)}{n_{mut} n_{wt}} \begin{bmatrix} n_{mut} & -n_{mut} \\ -n_{mut} & n_{mut} + n_{wt} \end{bmatrix}$$

Importantly, the estimate for the coefficient measuring the effect that the presence of the mutation has variance

$$\mathbb{V} \left[\beta_{s,mut}^{(OLS)} \right] = \frac{(\sigma_s^2 + \tau_s^2)(n_{mut} + n_{wt})}{n_{mut} n_{wt}}$$

When there is only 1 mutated sample, as per the motivation of this work, this reduces to

$$\mathbb{V} \left[\beta_{s,mut}^{(OLS)} \right] = \frac{(\sigma_s^2 + \tau_s^2)(1 + n_{wt})}{n_{wt}} \quad (\text{A.5})$$

A.3 Statistical moments of the James-Stein estimator

A.3.1 Expected value of the James-Stein estimator

We can use a Taylor expansion around B_1 to approximate the expected value of $\hat{B}_1^{(JS)}$. Consider:

$$\hat{B}_1^{(JS)} = \left(1 - \frac{c}{\left(\hat{B}_1^{(OLS)} \right)^T \Sigma^{-1} \hat{B}_1^{(OLS)}} \right) \hat{B}_1^{(OLS)}$$

where

$$\begin{aligned} \hat{B}_1^{(OLS)} &\sim N_{|S|}(B_1, \Sigma) \\ \Sigma_{s,t} &= \begin{cases} \left(\frac{n_{wt}+1}{n_{wt}} \right) (\sigma_s^2 + \tau_s^2) & s = t \\ 0 & s \neq t \end{cases} \end{aligned}$$

Let $u = \Sigma^{-1/2} \hat{B}_1^{(OLS)}$. Then

$$\begin{aligned} \mathbb{E} \left[\hat{B}_1^{(JS)} \right] &= \mathbb{E} \left[\hat{B}_1^{(OLS)} \right] - c \Sigma^{1/2} \mathbb{E} \left[\frac{u}{\|u\|^2} \right] \\ &= B_1 - c \Sigma^{1/2} \mathbb{E} \left[\frac{u}{\|u\|^2} \right] \Sigma^{1/2} \end{aligned}$$

Expanding $\frac{u}{\|u\|^2}$ around $a = \Sigma^{-1/2} B_1$ gives:

$$\begin{aligned} \mathbb{E} \left[\hat{B}_1^{(JS)} \right] &= B_1 - c \Sigma^{1/2} \mathbb{E} \left[\frac{a}{\|a\|^2} + \left(\frac{1}{\|a\|^2} - \frac{2}{\|a\|^4} a a^T \right) (u - a) + \mathcal{O}(\|u - a\|^2) \right] \\ &= \left(1 - \frac{c}{B_1^T \Sigma^{-1} B_1} \right) B_1 + \mathcal{O}(\|u - a\|^2) \end{aligned}$$

As long as the number of transcripts being considered, $|S|$, is not large, and that the true coefficient of variation is not large (i.e. that $\|u - a\|^2 \ll \|B_1\|^2$), the Taylor approximation is close to

$$\mathbb{E} [\hat{B}_1^{(JS)}] \approx \left(1 - \frac{c}{B_1^T \Sigma^{-1} B_1}\right) B_1 \quad (\text{A.6})$$

Thus the JS estimator is an estimate of B_1 that is biased towards 0.

A.3.2 Variance of the James-Stein estimator

The MSE of the JS estimator is related to its variance.

$$\mathbb{E} [\|\hat{B}_1^{(JS)} - B_1\|^2] = \sum_{s \in S} \mathbb{E} \left[\left(\hat{B}_{1,s}^{(JS)} - B_{1,s} \right)^2 \right] = \sum_{s \in S} \mathbb{V} [\hat{B}_{1,s}^{(JS)}]$$

By [REF 3], $\mathbb{E} [\|\hat{B}_1^{(JS)} - B_1\|^2] \leq \mathbb{E} [\|\hat{B}_1^{(OLS)} - B_1\|^2]$. However, this does not imply that $\mathbb{V} [\hat{B}_{1,s}^{(JS)}] \leq \mathbb{V} [\hat{B}_{1,s}^{(OLS)}] \forall s \in S$. Some transcripts may have larger variances than the OLS estimator, but all transcripts in aggregate will have a smaller MSE. This is still desirable if the goal is to find if there is an effect on any transcripts in the set S , instead of a particular one within the set.

To calculate the variance for each individual transcript, a similar approach with Taylor expansions can be used, as above.

$$\begin{aligned} & \mathbb{V} [\hat{B}_1^{(JS)}] \\ & \approx \mathbb{E} \left[\hat{B}_1^{(JS)} \left(\hat{B}_1^{(JS)} \right)^T \right] - \left(1 - \frac{c}{B_1^T \Sigma^{-1} B_1} \right)^2 B_1 B_1^T \\ & = \Sigma^{1/2} \mathbb{E} \left[uu^T - \frac{2c}{u^T u} uu^T + \left(\frac{c}{u^T u} \right)^2 uu^T \right] \Sigma^{1/2} - \left(1 - \frac{c}{B_1^T \Sigma^{-1} B_1} \right)^2 B_1 B_1^T \end{aligned}$$

where, again, $u = \Sigma^{-1/2} \hat{B}_1^{(OLS)}$. Expanding about $a = \Sigma^{-1/2} B_1$ yields:

$$\mathbb{V} [\hat{B}_1^{(JS)}] = \left(1 - \frac{2c}{B_1^T \Sigma^{-1} B_1} \right) \Sigma - \frac{2c}{(B_1^T \Sigma^{-1} B_1)^2} B_1 B_1^T + \mathcal{O}(\|u - a\|^4)$$

Under similar conditions of the number of transcripts under consideration, $|S|$, and $\|u - a\|^2$, we then have that

$$\mathbb{V}[\hat{\mathbf{B}}_1^{(JS)}] \approx \left(1 - \frac{2c}{\mathbf{B}_1^T \Sigma^{-1} \mathbf{B}_1}\right) \Sigma - \frac{2c}{(\mathbf{B}_1^T \Sigma^{-1} \mathbf{B}_1)^2} \mathbf{B}_1 \mathbf{B}_1^T \quad (\text{A.7})$$

Since the diagonal elements of $\frac{2c}{(\mathbf{B}_1^T \Sigma^{-1} \mathbf{B}_1)^2} \mathbf{B}_1 \mathbf{B}_1^T$ are all ≥ 0 and $0 \leq \left(1 - \frac{2c}{\mathbf{B}_1^T \Sigma^{-1} \mathbf{B}_1}\right) \leq 1 \forall c > 0$, the variance than of the JS estimators are smaller than the OLS estimators. The resulting Wald test statistics for the fold change coefficient of transcript s in the OLS and JS cases can be summarized as follows:

$$W_s^{(OLS)} = \frac{\left(\hat{\mathbf{B}}_{1,s}^{(OLS)}\right)^2}{\Sigma_{s,s}} \quad (\text{A.8})$$

$$W_s^{(JS)} = \frac{\left(1 - \frac{2c}{\left(\hat{\mathbf{B}}_1^{(OLS)}\right)^T \Sigma^{-1} \hat{\mathbf{B}}_1^{(OLS)}}\right)^2 \left(\hat{\mathbf{B}}_{1,s}^{(OLS)}\right)^2}{\left(1 - \frac{2c}{\left(\hat{\mathbf{B}}_1^{(OLS)}\right)^T \Sigma^{-1} \hat{\mathbf{B}}_1^{(OLS)}}\right) \Sigma_{s,s} - \frac{2c}{\left(\left(\hat{\mathbf{B}}_1^{(OLS)}\right)^T \Sigma^{-1} \hat{\mathbf{B}}_1^{(OLS)}\right)^2} \left(\hat{\mathbf{B}}_{1,s}^{(OLS)}\right)^2} \quad (\text{A.9})$$

The coefficient of $\hat{\mathbf{B}}_{1,s}^{(OLS)}$ in the numerator is larger than the coefficient of Σ in the denominator since $(1 - a)^2 = 1 - 2a + a^2 > 1 - 2a \forall a \in \mathbb{R}$. This implies that the Wald test statistics will be larger for the JS estimator than for the OLS estimator. Thus the JS method will produce more positive calls, in general, than the OLS method.

Notably, the variance of the JS estimator is a function of both the mean and variance of the transcripts under consideration. This is in contrast to the OLS estimator, which is solely a function of the variance. Additionally, the off-diagonal elements of the matrix $\mathbf{B}_1 \mathbf{B}_1^T$ imply that the JS fold change estimates are not independent of each other. This, again, contrasts with the OLS estimator, where the diagonal covariance matrix, Σ , implies that the fold change estimates are themselves independent of each other. The effect of this dependence on statistical inference is a function of the variance and true fold change, as can be seen from the $\frac{2c}{(\mathbf{B}_1^T \Sigma^{-1} \mathbf{B}_1)^2}$ coefficient. While rarely true in practice, this statistical dependence can affect the results of statistical inference, in theory. For most purposes, is not expected to have a large effect on the results of statistical inference.

Glossary

3C chromatin conformation capture

AR androgen receptor

ChIP-seq chromatin immunoprecipitation sequencing

CPC-GENE Canadian Prostate Cancer Genome Network

crRNA CRISPR RNA

CRE *cis*-regulatory element

DEPMAP Cancer Dependency Map

DHS DNase I hypersensitive sites

FN false negative

FP false positive

FOX forkhead box

GLM generalized linear model

gRNA guide RNA

IID independent and identically distributed

JS James-Stein

kbp kilobase

KO knockout

MSE mean square error

mCRPC metastatic castration-resistant prostate cancer

OLS ordinary least squares

mRNA messenger RNA

PCa prostate cancer

RNAi RNA interference

RNA-seq RNA sequencing

shRNA small hairpin RNA

siRNA small interfering RNA

SNV single nucleotide variants

SRA Sequence Read Archive

SV structural variant

TAD topologically associated domain

TCGA The Cancer Genome Atlas

TN true negative

TP true positive

TF transcription factor

tracrRNA trans-activating CRISPR RNA

UTR untranslated region

WGS whole genome sequencing

WT wild-type

References

1. Gierliński, M. *et al.* Statistical Models for RNA-Seq Data Derived from a Two-Condition 48-Replicate Experiment. en. *Bioinformatics* **31**, 3625–3630. ISSN: 1367-4803, 1460-2059 (Nov. 2015).
2. Stein, C. *Inadmissibility of the Usual Estimator for the Mean of a Multivariate Normal Distribution* en. in *Contribution to the Theory of Statistics* **3** (University of California Press, Berkeley, California, USA, Dec. 1956), 197–206. ISBN: 978-0-520-31388-0.
3. Bock, M. E. Minimax Estimators of the Mean of a Multivariate Normal Distribution. en. *The Annals of Statistics* **3**, 209–218. ISSN: 0090-5364 (Jan. 1975).
4. Pimentel, H., Bray, N. L., Puente, S., Melsted, P. & Pachter, L. Differential Analysis of RNA-Seq Incorporating Quantification Uncertainty. en. *Nature Methods* **14**, 687–690. ISSN: 1548-7105 (July 2017).
5. Yi, L., Pimentel, H., Bray, N. L. & Pachter, L. Gene-Level Differential Analysis at Transcript-Level Resolution. *Genome Biology* **19**, 53. ISSN: 1474-760X (Apr. 2018).
6. Simon Andrews. *FastQC: A Quality Control Tool for High Throughput Sequence Data* 2010.
7. Bray, N. L., Pimentel, H., Melsted, P. & Pachter, L. Near-Optimal Probabilistic RNA-Seq Quantification. en. *Nature Biotechnology* **34**, 525–527. ISSN: 1546-1696 (May 2016).
8. R Core Team. *R: A Language and Environment for Statistical Computing*. R Foundation for Statistical Computing. Vienna, Austria, 2013.
9. Benjamini, Y. & Hochberg, Y. Controlling the False Discovery Rate: A Practical and Powerful Approach to Multiple Testing. *Journal of the Royal Statistical Society. Series B (Methodological)* **57**, 289–300. ISSN: 0035-9246 (1995).